

# SQL SELECT III

G51DBS Database Systems

Jason Atkin

# Last Lecture

- More SQL SELECT
  - Aliases
  - 'Self-Joins'
  - Subqueries
  - IN, EXISTS, ANY, ALL
  - LIKE

# Example 2 from last week

Student

sID	sName	sAddress	sYear
1	Smith	5 Arnold Close	2
2	Brooks	7 Holly Avenue	2
3	Anderson	15 Main Street	3
4	Evans	Flat 1a, High Street	2
5	Harrison	Newark Hall	1
6	Jones	Southwell Hall	1

**Find a list of the names of any students who are enrolled on at least one module alongside 'Evans'**

Module

mCode	mCredits	mTitle
G51DBS	10	Database Systems
G51PRG	20	Programming
G51IAI	10	Artificial Intelligence
G52ADS	10	Algorithms

Enrolment

sID	mCode
1	G52ADS
2	G52ADS
5	G51DBS
5	G51PRG
5	G51IAI
4	G52ADS
6	G51PRG
6	G51IAI

# Student NATURAL JOIN Enrolment

Student NATURAL JOIN Enrolment

sID	sName	sAddress	sYear	mCode
1	Smith	5 Arnold Close	2	G52ADS
2	Brooks	7 Holly Avenue	2	G52ADS
<del>3</del>	<del>Anderson</del>	<del>15 Main Street</del>	<del>3</del>	
4	Evans	Flat 1a, High Street	2	G52ADS
5	Harrison	Newark Hall	1	G51DBS
5	Harrison	Newark Hall	1	G51PRG
5	Harrison	Newark Hall	1	G51IAI
6	Jones	Southwell Hall	1	G51PRG
6	Jones	Southwell Hall	1	G51IAI

**Question did not say we could not use a join**

## Example 2: 2 joins and a sub-query

- Find a list of the names of any students who are enrolled on at least one module alongside 'Evans'

```
SELECT DISTINCT S1.sName, E1.mCode
FROM
Student S1 NATURAL JOIN Enrolment E1
WHERE E1.mCode IN
( SELECT E2.mCode FROM
Enrolment E2 NATURAL JOIN Student S2
WHERE S2.sName = 'Evans' );
```

Could also join the two sub-queries on `E1.mCode = E2.mCode`

## Example 2: Three joins

- Find a list of the names of any students who are enrolled on at least one module alongside 'Evans'

```
SELECT DISTINCT S1.sName, E1.mCode
FROM
(Student S1 NATURAL JOIN Enrolment E1)
INNER JOIN
(Student S2 NATURAL JOIN Enrolment E2)
USING (mCode)
WHERE S2.sName = 'Evans';
```

Cannot NATURAL JOIN the parts here! (sName, mCode etc would match)

## Example 2: Sub queries, get names

```
SELECT DISTINCT S1.sName FROM Student S1
WHERE S1.sID IN
  ( SELECT E1.sID FROM Enrolment E1
    WHERE E1.mCode IN
      ( SELECT E2.mCode FROM Enrolment E2
        WHERE E2.sID IN
          ( SELECT S2.sID FROM Student S2
            WHERE S2.sName = 'Evans') ) );
```

As long as we only want the student name, not the module id (mCode not in table)

# Coursework Reminder

- Don't forget the coursework
  - Create an ER diagram
  - Create the SQL to create the tables
  - Create the SQL to populate the tables
- **Deadline: 4pm Thursday 13th March**
- Hopefully you have had/will have had a tutorial with your tutor by then



# This Lecture

- More SQL SELECT
  - ORDER BY
  - Aggregate functions
  - GROUP BY and HAVING
  - UNION
- Further reading
  - The Manga Guide to Databases, Chapter 4
  - Database Systems, Chapter 6

# SQL SELECT Overview

**SELECT**

**[DISTINCT | ALL] <column-list>**

**FROM <table-names>**

**[WHERE <condition>]**

**[GROUP BY <column-list>]**

**[HAVING <condition>]**

**[ORDER BY <column-list>]**

*([] optional, | or)*

**ORDER BY**

# ORDER BY

- The ORDER BY clause sorts the results of a query
  - You can sort in ascending or descending order
  - Defaults to ascending if neither is specified
  - Multiple columns can be given
  - You cannot order by a column which isn't in the result

```
SELECT <columns>
      FROM <tables>
      WHERE <condition>
      ORDER BY <cols>
              [ASC | DESC]
```

# ORDER BY

**SELECT \* FROM Grades  
ORDER BY Mark**

Grades

Name	Code	Mark
John	DBS	56
John	IAI	72
Mary	DBS	60
James	PR1	43
James	PR2	35
Jane	IAI	54

Name	Code	Mark
James	PR2	35
James	PR1	43
Jane	IAI	54
John	DBS	56
Mary	DBS	60
John	IAI	72

# ORDER BY

```
SELECT * FROM Grades  
ORDER BY Code ASC,  
Mark DESC
```

Grades

Name	Code	Mark
John	DBS	56
John	IAI	72
Mary	DBS	60
James	PR1	43
James	PR2	35
Jane	IAI	54

Name	Code	Mark
Mary	DBS	60
John	DBS	56
John	IAI	72
Jane	IAI	54
James	PR1	43
James	PR2	35

# Constants and Arithmetic

# Constants and Arithmetic

- As well as columns, a SELECT statement can also be used to
  - Select constants
  - Compute arithmetic expressions
  - Evaluate functions
- Often helpful to use an alias when dealing with expressions or functions

```
SELECT Mark / 100  
FROM Grades
```

```
SELECT Salary + Bonus  
FROM Employee
```

```
SELECT 1.175 * Price  
AS 'Price inc. VAT'  
FROM Products
```



# Aggregate Functions

- Aggregate functions compute summaries of data in a table
  - Most aggregate functions work on a single column of numerical data
  - **COUNT** ( \* ) works on the entire row
- It is good to use an alias to name the result
- Aggregate functions
  - **COUNT**: The number of rows
  - **SUM**: The sum of the entries in the column
  - **AVG**: The average entry in a column
  - **MIN**, **MAX**: The minimum and maximum entries in a column

# COUNT : guess the results...

Grades

Name	Code	Mark
John	DBS	56
John	IAI	72
Mary	DBS	60
James	PR1	43
James	PR2	35
Jane	IAI	56
Jane	PR1	NULL

```
SELECT
    COUNT(*) AS Count
FROM Grades
```

Count
?

```
SELECT
    COUNT(Mark)
    AS Count
FROM Grades
```

Count
?

```
SELECT
    COUNT(DISTINCT Mark)
    AS Count
FROM Grades
```

Count
?

# COUNT

Grades

Name	Code	Mark
John	DBS	56
John	IAI	72
Mary	DBS	60
James	PR1	43
James	PR2	35
Jane	IAI	56
Jane	PR1	NULL

**SELECT**

**COUNT(\*) AS Count  
FROM Grades**

**Count**

7

7 rows

**SELECT**

**COUNT(Mark)  
AS Count  
FROM Grades**

**Count**

6

6 non-NULL marks

**SELECT**

**COUNT(DISTINCT Mark)  
AS Count  
FROM Grades**

**Count**

5

5 different  
non-NULL marks

# SUM, MIN/MAX and AVG

Grades

Name	Code	Mark
John	DBS	56
John	IAI	72
Mary	DBS	60
James	PR1	43
James	PR2	35
Jane	IAI	54

**SELECT**

**SUM(Mark) AS Total  
FROM Grades**

**Total**

320

**SELECT**

**MAX(Mark) AS Best  
FROM Grades**

**Best**

72

**SELECT**

**AVG(Mark) AS Mean  
FROM Grades**

**Mean**

53.33

Note: Work on ONE column. Ignore NULL values.

# Aggregate Functions

- You can combine aggregate functions using arithmetic

**SELECT**

**MAX(Mark) - MIN(Mark)**

**AS Range**

**FROM Grades**

Grades

Name	Code	Mark
John	DBS	56
John	IAI	72
Mary	DBS	60
James	PR1	43
James	PR2	35
Jane	IAI	54

**MAX(Mark) = 72**

**MIN(Mark) = 35**

**Range**

**37**

# Example

Modules

Code	Title	Credits
DBS	Database Systems	10
GRP	Group Project	20

- Find John's average mark, weighted by the credits of each module

Grades

Name	Code	Mark
John	DBS	56
John	IAI	72

**SELECT**

???

**FROM** ???

**WHERE** ???

# Example

Modules

Code	Title	Credits
DBS	Database Systems	10
GRP	Group Project	20

- Find John's average mark, weighted by the credits of each module

Grades

Name	Code	Mark
John	DBS	56
John	IAI	72

**SELECT**

???

**FROM Modules, Grades**

**WHERE Modules.Code = Grades.Code  
AND Grades.Name = 'John'**

# Example

Modules

Code	Title	Credits
DBS	Database Systems	10
GRP	Group Project	20

- Find John's average mark, weighted by the credits of each module

Grades

Name	Code	Mark
John	DBS	56
John	IAI	72

```
SELECT
SUM(Mark*Credits) / SUM(Credits)
  AS 'Final Mark'
FROM Modules, Grades

WHERE Modules.Code = Grades.Code
      AND Grades.Name = 'John'
```



Group by

# GROUP BY

- Sometimes we want to apply aggregate functions to groups of rows
  - The GROUP BY clause achieves this
- ```
SELECT <cols1>  
FROM <tables>  
GROUP BY <cols2>
```
- Example, find the average mark of each student individually rather than all together

# GROUP BY

```
SELECT <cols1>  
FROM <tables>  
GROUP BY <cols2>
```

- Every entry in <cols1> should be
  - in <cols2>
  - **or** be a constant
  - **or** be an aggregate function
    - ↑  
i.e. how to combine  
The rows returned

- You can have **WHERE** and **ORDER BY** clauses as well as a **GROUP BY** clause

1. Tables are joined
2. **WHERE** clauses remove rows
3. **GROUP BY** clauses combine remaining rows together

# GROUP BY

**SELECT**

**Name,**

**AVG(Mark) AS Average**

**FROM Grades**

**GROUP BY Name**

Grades

| Name  | Code | Mark |
|-------|------|------|
| John  | DBS  | 56   |
| John  | IAI  | 72   |
| Mary  | DBS  | 60   |
| James | PR1  | 43   |
| James | PR2  | 35   |
| Jane  | IAI  | 54   |

| Name  | Average |
|-------|---------|
| John  | 64      |
| Mary  | 60      |
| James | 39      |
| Jane  | 54      |

# GROUP BY

Sales

| Month | Department | Value |
|-------|------------|-------|
| March | Fiction    | 20    |
| March | Travel     | 30    |
| March | Technical  | 40    |
| April | Fiction    | 10    |
| April | Fiction    | 30    |
| April | Travel     | 25    |
| April | Fiction    | 20    |
| May   | Fiction    | 20    |
| May   | Travel     | 50    |

- Find the total value of the sales for each department in each month
  - Can group by Month then Department or Department then Month
  - Same results, but **possibly** produced in a different order

# GROUP BY

```
SELECT Month, Department,  
       SUM (Value) AS Total  
FROM Sales  
GROUP BY Month, Department
```

| Month | Department | Total |
|-------|------------|-------|
| April | Fiction    | 60    |
| April | Travel     | 25    |
| March | Fiction    | 20    |
| March | Technical  | 40    |
| March | Travel     | 30    |
| May   | Fiction    | 20    |
| May   | Technical  | 50    |

```
SELECT Month, Department,  
       SUM (Value) AS Total  
FROM Sales  
GROUP BY Department, Month
```

| Month | Department | Total |
|-------|------------|-------|
| April | Fiction    | 60    |
| March | Fiction    | 20    |
| May   | Fiction    | 20    |
| March | Technical  | 40    |
| May   | Technical  | 50    |
| April | Travel     | 25    |
| March | Travel     | 30    |

# GROUP BY Rules

- GROUP BY works slightly differently in MySQL than in other DBMSs.
- **Usually**, every column you name in your output must either be one of the GROUP BY columns or an aggregate function...
- For example:

```
SELECT  
    ID, Name, AVG(Mark)  
FROM Grades  
GROUP BY ID, Name
```

Grades

| Name  | ID | Code | Mark |
|-------|----|------|------|
| John  | 1  | DBS  | 56   |
| John  | 1  | IAI  | 72   |
| Mary  | 2  | DBS  | 60   |
| James | 3  | PR1  | 43   |
| James | 3  | PR2  | 35   |
| Jane  | 4  | IAI  | 54   |
| Jane  | 4  | DBS  | 32   |
| Jane  | 4  | PR1  | 73   |
| Bill  | 5  | DBS  | 54   |
| Bill  | 5  | IAI  | 63   |
| Bill  | 5  | PR1  | 12   |

# GROUP BY Rules

- In MySQL, *for convenience*, you are **allowed** to break this rule.
- You do not have to select a **SELECT** a column even if it's in your **GROUP BY** clause
- You can **SELECT** columns which are not in your **GROUP BY** clause
- **Despite this, you should follow the ISO standard where possible**
  - Avoids problems if you use a different DBMS in the future
- This can lead to peculiar output when multiple values are options for one row
- If there were multiple options for the same row, MySQL would be free to choose a random one for each name
  - “The server is free to choose any value from each group, so unless they are the same, the values chosen are indeterminate.” **Source:** **MySQL online documentation**



**HAVING**

# HAVING

- HAVING is like a WHERE clause, except that it only applies to the **results** of a GROUP BY query
- It can be used to **select groups** which satisfy a given condition

```
SELECT Name,  
       AVG(Mark) AS Average  
FROM Grades  
GROUP BY Name  
HAVING AVG(Mark) >= 40
```

| Name | Average |
|------|---------|
| John | 64      |
| Mary | 60      |
| Jane | 54      |

# WHERE and HAVING

- **WHERE refers to the rows** of tables, so cannot make use of aggregate functions
- **HAVING refers to the groups** of rows, and so cannot use columns which are not in the **GROUP BY** or an aggregate function
- Think of a query being processed as follows:
  1. Tables are joined
  2. **WHERE** clauses
  3. **GROUP BY** clauses and aggregates
  4. Column selection
  5. **HAVING** clauses
  6. **ORDER BY**

UNION

# UNION

- UNION, INTERSECT and EXCEPT
  - These treat the tables as sets and are the usual set operators of union, intersection and difference
  - We'll be concentrating on UNION
- They all combine the results from two select statements
- The results of the two selects should have the **same columns and data types**

# UNION

Grades

| Name  | Code | Mark |
|-------|------|------|
| Jane  | IAI  | 52   |
| John  | DBS  | 56   |
| John  | IAI  | 72   |
| James | PR1  | 43   |
| James | PR2  | 35   |
| Mary  | DBS  | 60   |

- **Question:**  
Find, in a single query, the average mark for each student and the average mark overall

# UNION

- The average for each student:

```
SELECT Name, AVG(Mark) AS Average  
FROM Grades GROUP BY Name
```

- The average overall

```
SELECT 'Total' AS Name, AVG(Mark) AS Average  
FROM Grades
```

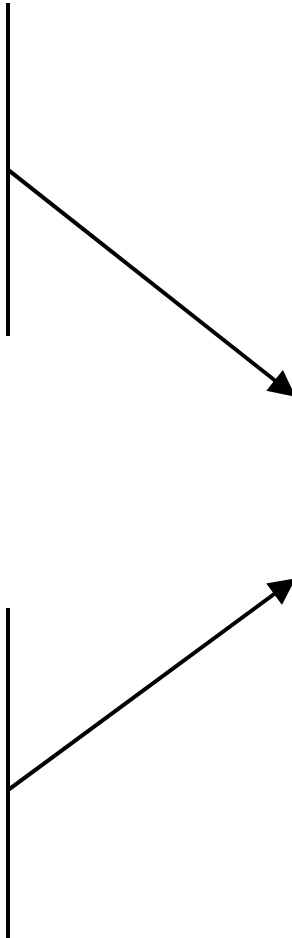
- Note - this has the **same columns** as average by student, i.e. 'Name' and 'Average'

# UNION

```
SELECT Name,  
       AVG(Mark) AS Average  
FROM Grades  
GROUP BY Name
```

UNION

```
SELECT  
  'Total' AS Name,  
  AVG(Mark) AS Average  
FROM Grades
```



| Name  | Average |
|-------|---------|
| Jane  | 52      |
| John  | 64      |
| James | 39      |
| Mary  | 60      |
| Total | 53      |



Combining these things...

# The Final SELECT Example

- Examiners' reports
  - We want a list of students and their average mark
  - For first and second years the average is for that year
  - For finalists it is 40% of the second year plus 60% of the final year average
- We want the results
  - Sorted by year, then average mark (high to low) then last name, first name and finally ID
  - To take into account the number of credits each module is worth
  - To be produced by a single query

# Tables for the Example

Student

| ID | First | Last | Year |
|----|-------|------|------|
|----|-------|------|------|

Grade

| ID | Code | Mark | YearTaken |
|----|------|------|-----------|
|----|------|------|-----------|

Module

| Code | Title | Credits |
|------|-------|---------|
|------|-------|---------|

# Getting Started

- Finalists should be treated differently to other years
  - Write one SELECT for the finalists
  - Write a second SELECT for the first and second years
  - Join the results using a UNION

<QUERY FOR FINALISTS>

UNION

<QUERY FOR OTHERS>

# Table Joins

- Both subqueries need information from all of the tables:
  - The student ID, name and year
  - The marks for each module and the year taken
  - The number of credits for each module
- This is an obvious natural join operation
  - Because we're practicing, we're going to use a standard CROSS JOIN and WHERE clause

# The Query So Far

```
SELECT <some information>
  FROM Student, Module, Grade
 WHERE Student.ID = Grade.ID
    AND Module.Code = Grade.Code
    AND <student is in third year>
```

UNION

```
SELECT <some information>
  FROM Student, Module, Grade
 WHERE Student.ID = Grade.ID
    AND Module.Code = Grade.Code
    AND <student is in first or second year>
```

# Information for Finalists

- We must retrieve
  - Computed average mark, weighted 40-60 across years 2 and 3
  - First year marks must be ignored
  - The ID, Name and Year are needed as they are used for ordering
- The average is difficult
  - We don't have any statements to separate years 2 and 3 easily
  - We can exploit the fact that  $40 = 20 * 2$  and  $60 = 20 * 3$ , so **YearTaken and the weighting have the same relationship**

# Information for Finalists

```
SELECT Year, Student.ID, Last, First,  
       SUM((20*YearTaken)/100)*Mark*Credits)/120  
       AS AverageMark  
FROM Student, Module, Grade  
WHERE Student.ID = Grade.ID  
      AND Module.Code = Grade.Code  
      AND YearTaken IN (2,3)  
      AND Year = 3  
GROUP BY Year, Student.ID, First, Last
```



# Information for Others

- Other students are easier than finalists
  - We just need their average marks where YearTaken and Year are the same
  - As before, we need ID, Name and Year for ordering

# Information for Others

```
SELECT Year, Student.ID, Last, First,  
       SUM(Mark*Credits)/120 AS AverageMark  
FROM Student, Module, Grade  
WHERE Student.ID = Grade.ID  
      AND Module.Code = Grade.Code  
      AND YearTaken = Year  
      AND Year IN (1,2)  
GROUP BY Year, Student.ID, First, Last
```

# The Final Query

```
SELECT Year, Student.ID, Last, First,  
       SUM((20*YearTaken)/100)*Mark*Credits)/120 AS AverageMark  
FROM Student, Module, Grade  
WHERE Student.ID = Grade.ID AND Module.Code = Grade.Code  
      AND YearTaken IN (2,3)      AND Year = 3  
GROUP BY Year, Student.ID, Last, First
```

UNION

```
SELECT Year, Student.ID, Last, First, SUM(Mark*Credits)/120 AS  
      AverageMark  
FROM Student, Module, Grade  
WHERE Student.ID = Grade.ID AND Module.Code = Grade.Code  
      AND YearTaken = Year      AND Year IN (1,2)  
GROUP BY Year, Student.ID, Last, First
```

```
ORDER BY Year desc, AverageMark desc, Last, First, ID
```

# Example Output

| Year | Student.ID | Last     | First   | AverageMark |
|------|------------|----------|---------|-------------|
| 3    | 11014456   | Andrews  | John    | 81          |
| 3    | 11013891   | Smith    | Mary    | 78          |
| 3    | 11014012   | Jones    | Steven  | 76          |
| 3    | 11013204   | Brown    | Amy     | 76          |
| 3    | 11014919   | Robinson | Paul    | 74          |
| 2    | 11014584   | Edwards  | Robert  | 73          |
| ...  |            |          |         |             |
| 1    | 11027871   | Green    | Michael | 45          |
| 1    | 11024298   | Hall     | David   | 43          |
| 1    | 11024826   | Wood     | James   | 40          |
| 1    | 11027621   | Clarke   | Stewart | 39          |
| 1    | 11024978   | Wilson   | Sarah   | 36          |
| 1    | 11026563   | Taylor   | Matthew | 34          |
| 1    | 11027625   | Williams | Paul    | 31          |

# Next Lecture

- Missing Information
  - Nulls and the Relational Model
  - Outer Joins
  - Default Values
- Further reading
  - The Manga Guide to Databases, Chapter 2
  - Database Systems, Chapter 4

# Missing Information

G51DBS Database Systems

Jason Atkin

[jaa@cs.nott.ac.uk](mailto:jaa@cs.nott.ac.uk)

# This Lecture

- Missing Information
  - Nulls and the Relational Model
  - Outer Joins
  - Default Values
- Further reading
  - The Manga Guide to Databases, Chapter 2
  - Database Systems, Chapter 4

# Missing Information

- Sometimes we don't know what value an entry in a relation should have
  - We know that there is a value, but don't know what it is
  - There is no value at all that makes any sense
- Two main methods have been proposed to deal with this
  - NULLs can be used as markers to show that information is missing
  - A default value can be used to represent the missing value



# NULLs

- NULL is a placeholder for missing or unknown value of an attribute. It is not itself a value.
- Codd proposed to distinguish two kinds of NULLs:
  - A-marks: data **Applicable** but not known (for example, someone's age)
  - I-marks: data is **Inapplicable** (telephone number for someone who does not have a telephone, or spouse's name for someone who is not married)

# Problems with NULLs

- Problems with extending relational algebra operations to NULLs:
  - Defining selection operation:  
if we check tuples for some property like  $\text{Mark} > 40$  and for some tuple, Mark is NULL, do we include it?
  - Comparing tuples in two relations: are the two tuples  $\langle \text{John}, \text{NULL} \rangle$  and  $\langle \text{John}, \text{NULL} \rangle$  the same or not?
- Additional problems for SQL:
  - Do we treat NULLs as duplicates?
  - Do we include them in count, sum, average and if yes, how?
  - How do arithmetic operations behave when an argument is NULL?

# Theoretical Solutions

- Use three-valued logic instead of classical two-valued logic to evaluate conditions
- When there are no NULLs around, conditions evaluate to true or false, but if a null is involved, a condition might evaluate to the third value ('undefined', or 'unknown')
- This is the idea behind testing conditions in the WHERE clause of SQL SELECT: only tuples where the condition evaluates to true are returned
  - i.e. NOT 'false' **OR** 'undefined'

# 3-valued logic

- Results if condition involves a boolean combination:

| a       | b       | a OR b  | a AND b | a == b  |
|---------|---------|---------|---------|---------|
| True    | True    | True    | True    | True    |
| True    | False   | True    | False   | False   |
| False   | True    | True    | False   | False   |
| False   | False   | False   | False   | True    |
| True    | Unknown | True    | Unknown | Unknown |
| Unknown | True    | True    | Unknown | Unknown |
| False   | Unknown | Unknown | False   | Unknown |
| Unknown | False   | Unknown | False   | Unknown |
| Unknown | Unknown | Unknown | Unknown | Unknown |

# SQL NULLs in Conditions

```
SELECT *  
FROM Employee  
Where  
    salary > 15,000;
```

- **salary > 15,000**  
evaluates to 'unknown'  
on the last tuple
- so not included

Employee

| Name  | Salary |
|-------|--------|
| John  | 25,000 |
| Mark  | 15,000 |
| Anne  | 20,000 |
| Chris | NULL   |

| Name | Salary |
|------|--------|
| John | 25,000 |
| Anne | 20,000 |

# SQL NULLs in Conditions

```
SELECT *  
FROM Employee  
Where  
    Salary > 15,000  
    OR Name = 'Chris';
```

Employee

| Name  | Salary |
|-------|--------|
| John  | 25,000 |
| Mark  | 15,000 |
| Anne  | 20,000 |
| Chris | NULL   |

- **Salary > 15,000 OR Name = 'Chris' is Unknown OR TRUE on the last tuple**
- **Unknown OR TRUE = TRUE**

| Name  | Salary |
|-------|--------|
| John  | 25,000 |
| Anne  | 20,000 |
| Chris | NULL   |

# SQL NULLs in Arithmetic

```
SELECT  
  Name,  
  Salary * 0.05 AS Bonus  
FROM Employee;
```

- Arithmetic operations applied to NULLs result in NULLS

Employee

| Name  | Salary |
|-------|--------|
| John  | 25,000 |
| Mark  | 15,000 |
| Anne  | 20,000 |
| Chris | NULL   |

| Name  | Bonus |
|-------|-------|
| John  | 1,250 |
| Mark  | 750   |
| Anne  | 1,000 |
| Chris | NULL  |

# SQL NULLs in Aggregation

```
SELECT
  AVG(Salary) AS Average,
  COUNT(Salary) AS Count,
  SUM(Salary) AS Sum
FROM Employee;
```

Employee

| Name  | Salary |
|-------|--------|
| John  | 25,000 |
| Mark  | 15,000 |
| Anne  | 20,000 |
| Chris | NULL   |

- Average = 20,000
- Count = 3
- Sum = 60,000
- Using COUNT(\*) would give 4



# SQL NULLs in GROUP BY

```
SELECT  
    Salary,  
    COUNT(Name) AS Count  
FROM Employee  
GROUP BY Salary;
```

- NULLs are treated as equivalents in GROUP BY clauses

Employee

| Name  | Salary |
|-------|--------|
| John  | 25,000 |
| Mark  | 15,000 |
| Anne  | 20,000 |
| Jack  | NULL   |
| Sam   | 20,000 |
| Chris | NULL   |

| Salary | Count |
|--------|-------|
| NULL   | 2     |
| 15,000 | 1     |
| 20,000 | 2     |
| 25,000 | 1     |

# Outer Joins

# Outer Joins

- When we take the **join** of two relations we match up tuples which share values
- Some tuples have no match, and are 'lost' with an inner join
- These are called 'dangles'
- Joins do not normally include dangles
- Outer joins **include** dangles in the result and use NULLs to fill in the blanks
  - LEFT OUTER JOIN
  - RIGHT OUTER JOIN
  - FULL OUTER JOIN
- Outer Joins use ON like INNER JOIN
- Note: Cannot use 'USING' version

# Example: Inner Join

Student

| ID         | Name        |
|------------|-------------|
| 123        | John        |
| 124        | Mary        |
| 125        | Mark        |
| <b>126</b> | <b>Jane</b> |

Enrolment

| ID         | Code       | Mark      |
|------------|------------|-----------|
| 123        | DBS        | 60        |
| 124        | PRG        | 70        |
| 125        | DBS        | 50        |
| <b>128</b> | <b>DBS</b> | <b>80</b> |

← Dangles

Student INNER JOIN Enrolment ON Student.ID = Enrolment.ID

| ID  | Name | ID  | Code | Mark |
|-----|------|-----|------|------|
| 123 | John | 123 | DBS  | 60   |
| 124 | Mary | 124 | PRG  | 70   |
| 125 | Mark | 125 | DBS  | 50   |

# Outer Join Syntax

```
SELECT <cols>  
    FROM <table1> <type> OUTER JOIN <table2>  
        ON <condition>
```

Where <type> is one of **LEFT**, **RIGHT** or **FULL**

Example:

```
SELECT *  
    FROM Student LEFT OUTER JOIN Enrolment  
        ON Student.ID = Enrolment.ID
```

# Example: Left Outer Join

Student

| ID         | Name        |
|------------|-------------|
| 123        | John        |
| 124        | Mary        |
| 125        | Mark        |
| <b>126</b> | <b>Jane</b> |

Enrolment

| ID         | Code       | Mark      |
|------------|------------|-----------|
| 123        | DBS        | 60        |
| 124        | PRG        | 70        |
| 125        | DBS        | 50        |
| <b>128</b> | <b>DBS</b> | <b>80</b> |

← Dangles

Student LEFT OUTER JOIN Enrolment ON ...

| ID  | Name | ID   | Code | Mark |
|-----|------|------|------|------|
| 123 | John | 123  | DBS  | 60   |
| 124 | Mary | 124  | PRG  | 70   |
| 125 | Mark | 125  | DBS  | 50   |
| 126 | Jane | NULL | NULL | NULL |

# Example: Right Outer Join

Student

| ID  | Name |
|-----|------|
| 123 | John |
| 124 | Mary |
| 125 | Mark |
| 126 | Jane |

Enrolment

| ID  | Code | Mark |
|-----|------|------|
| 123 | DBS  | 60   |
| 124 | PRG  | 70   |
| 125 | DBS  | 50   |
| 128 | DBS  | 80   |

← Dangles

Student RIGHT OUTER JOIN Enrolment ON ...

| ID   | Name | ID  | Code | Mark |
|------|------|-----|------|------|
| 123  | John | 123 | DBS  | 60   |
| 124  | Mary | 124 | PRG  | 70   |
| 125  | Mark | 125 | DBS  | 50   |
| NULL | NULL | 128 | DBS  | 80   |

# Example: Full Outer Join

Student

| ID         | Name        |
|------------|-------------|
| 123        | John        |
| 124        | Mary        |
| 125        | Mark        |
| <b>126</b> | <b>Jane</b> |

Enrolment

| ID         | Code       | Mark      |
|------------|------------|-----------|
| 123        | DBS        | 60        |
| 124        | PRG        | 70        |
| 125        | DBS        | 50        |
| <b>128</b> | <b>DBS</b> | <b>80</b> |

← Dangles

Student FULL OUTER JOIN Enrolment ON ...

| ID   | Name | ID   | Code | Mark |
|------|------|------|------|------|
| 123  | John | 123  | DBS  | 60   |
| 124  | Mary | 124  | PRG  | 70   |
| 125  | Mark | 125  | DBS  | 50   |
| 126  | Jane | NULL | NULL | NULL |
| NULL | NULL | 128  | DBS  | 80   |



# Full Outer Join in MySQL

- Only Left and Right outer joins are supported in MySQL. If you really want a FULL outer join:

```
SELECT *  
  FROM Student FULL OUTER JOIN Enrolment  
    ON Student.ID = Enrolment.ID;
```

- Can be achieved using:

```
SELECT * FROM Student LEFT OUTER JOIN  
  Enrolment ON Student.ID = Enrolment.ID  
UNION  
SELECT * FROM Student RIGHT OUTER JOIN  
  Enrolment ON Student.ID = Enrolment.ID;
```

# Example

- Sometimes an outer join is the most practical approach. We may encounter NULL values, but may still wish to see the existing information
- Example: For students graduating in absentia, find a list of all student IDs, names, addresses, phone numbers and their final degree classifications.

# Example

Student

| ID  | Name | aID  | pID  | Grad |
|-----|------|------|------|------|
| 123 | John | 12   | 22   | C    |
| 124 | Mary | 23   | 90   | A    |
| 125 | Mark | 19   | NULL | A    |
| 126 | Jane | 14   | 17   | C    |
| 127 | Sam  | NULL | 101  | A    |

Phone

| pID | pNumber | pMobile     |
|-----|---------|-------------|
| 17  | 1111111 | 07856232411 |
| 22  | 2222222 | 07843223421 |
| 90  | 3333333 | 07155338654 |
| 101 | 4444444 | 07213559864 |

Degree

| ID  | Classification |
|-----|----------------|
| 123 | 1              |
| 124 | 2:1            |
| 125 | 2:2            |
| 126 | 2:1            |
| 127 | 3              |

Address

| aID | aStreet        | aTown      | aPostcode |
|-----|----------------|------------|-----------|
| 12  | 5 Arnold Close | Nottingham | NG12 1DD  |
| 14  | 17 Derby Road  | Nottingham | NG7 4FG   |
| 19  | 1 Main Street  | Derby      | DE1 5FS   |
| 23  | 7 Holly Avenue | Nottingham | NG6 7AR   |

# Example: INNER JOINS

- An Inner Join with Student and Address will ignore Student 127, who doesn't have an address record

Student

| ID  | Name | aID  | pID  | Grad |
|-----|------|------|------|------|
| 123 | John | 12   | 22   | C    |
| 124 | Mary | 23   | 90   | A    |
| 125 | Mark | 19   | NULL | A    |
| 126 | Jane | 14   | 17   | C    |
| 127 | Sam  | NULL | 101  | A    |

Address

| aID | aStreet        | aTown      | aPostcode |
|-----|----------------|------------|-----------|
| 12  | 5 Arnold Close | Nottingham | NG12 1DD  |
| 14  | 17 Derby Road  | Nottingham | NG7 4FG   |
| 19  | 1 Main Street  | Derby      | DE1 5FS   |
| 23  | 7 Holly Avenue | Nottingham | NG6 7AR   |

# Example: INNER JOINS

- An Inner Join with Student and Phone will ignore student 125, who doesn't have a phone record

Student

| ID  | Name | aID  | pID  | Grad |
|-----|------|------|------|------|
| 123 | John | 12   | 22   | C    |
| 124 | Mary | 23   | 90   | A    |
| 125 | Mark | 19   | NULL | A    |
| 126 | Jane | 14   | 17   | C    |
| 127 | Sam  | NULL | 101  | A    |

Phone

| pID | pNumber | pMobile     |
|-----|---------|-------------|
| 17  | 1111111 | 07856232411 |
| 22  | 2222222 | 07843223421 |
| 90  | 3333333 | 07155338654 |
| 101 | 4444444 | 07213559864 |

# Example

```
SELECT ID, Name, aStreet, aTown, aPostcode, pNumber,
       Classification
FROM Student LEFT OUTER JOIN Phone
      ON Student.pID = Phone.pID
LEFT OUTER JOIN Address
      ON Student.aID = Address.aID
INNER JOIN Degree ON Student.ID = Degree.ID
WHERE Grad = 'A';
```

Student

| ID | Name | aID | pID | Grad |
|----|------|-----|-----|------|
|----|------|-----|-----|------|

Phone

| pID | pNumber | pMobile |
|-----|---------|---------|
|-----|---------|---------|

Degree

| ID | Classification |
|----|----------------|
|----|----------------|

Address

| aID | aStreet | aTown | aPostcode |
|-----|---------|-------|-----------|
|-----|---------|-------|-----------|

# Example

| ID  | Name | aStreet        | aTown       | aPostcode   | pNumber     | Classification |
|-----|------|----------------|-------------|-------------|-------------|----------------|
| 124 | Mary | 7 Holly Avenue | Nottingham  | NG6 7AR     | 3333333     | 2:1            |
| 125 | Mark | 1 Main Street  | Derby       | DE1 5FS     | <b>NULL</b> | 2:2            |
| 127 | Sam  | <b>NULL</b>    | <b>NULL</b> | <b>NULL</b> | 4444444     | 3              |

- The records for students 125 and 127 have been preserved despite missing information

# Default Values



# Default Values

- Default values are an alternative to the use of NULLs
- If a value is not known then a particular **placeholder value** (the default) is used
- These are actual values, so don't need 3-value logic (3VL) etc.
- Default values can have more meaning than NULLs
  - 'none'
  - 'unknown'
  - 'not supplied'
  - 'not applicable'
- Not all defaults represent missing information. It depends on the situation

# Default Value Example

Parts

| ID | Name           | Weight | Quantity |
|----|----------------|--------|----------|
| 1  | Nut            | 10     | 20       |
| 2  | Bolt           | 15     | -1       |
| 3  | Nail           | 3      | 100      |
| 4  | Pin            | -1     | 30       |
| 5  | <b>Unknown</b> | 20     | 20       |
| 6  | Screw          | -1     | -1       |
| 7  | Brace          | 150    | 0        |

- Default values are
  - “Unknown” for Name
  - -1 for Weight and Quantity
- -1 is used for Wgt and Qty as it is not a sensible value, so cannot appear by accident as a real value
- There are still problems:  
`UPDATE Parts`  
`SET Quantity =`  
`Quantity + 5`

# Problems With Default Values

- Since defaults are real values:
- They can be updated like any other value
- You need to use a value that won't appear in any other circumstances
- They might not be interpreted properly
- Also, within SQL defaults must be of the same type as the column
- You can't have a string such as 'unknown' in a column of integers

# Splitting Tables

# Splitting Tables

- NULLs and defaults both try to fill entries with missing data
  - NULLs mark the data as missing
  - Defaults give some indication as to what sort of missing information we are dealing with
  - Defaults rely on having an 'invalid' value to use
- Often you can remove entries that have missing data
  - You can split the table up so that columns which might have NULLs are in separate tables
  - Entries that would be NULL are not present in these tables

# Splitting Tables Example

Parts

| ID | Name  | Weight | Quantity |
|----|-------|--------|----------|
| 1  | Nut   | 10     | 20       |
| 2  | Bolt  | 15     | NULL     |
| 3  | Nail  | 3      | 100      |
| 4  | Pin   | NULL   | 30       |
| 5  | NULL  | 20     | 20       |
| 6  | Screw | NULL   | NULL     |
| 7  | Brace | 150    | 0        |



Names

| ID | Name  |
|----|-------|
| 1  | Nut   |
| 2  | Bolt  |
| 3  | Nail  |
| 4  | Pin   |
| 6  | Screw |
| 7  | Brace |

Weights

| ID | Weight |
|----|--------|
| 1  | 10     |
| 2  | 15     |
| 3  | 3      |
| 5  | 20     |
| 7  | 150    |

Quantities

| ID | Name |
|----|------|
| 1  | 20   |
| 3  | 100  |
| 4  | 30   |
| 5  | 20   |
| 7  | 0    |

# Problems with Splitting Tables

- Splitting tables has other problems
  - Could introduce many more tables
  - Information gets spread out over the database
  - Queries become more complex and require many joins
- We can recover the original table, but
  - Requires Outer Joins
  - Reintroduces the NULL values, which means we're back to the original problem

SQL support for these things



# SQL Support

- SQL allows both NULLs and defaults:
- A table to hold data on employees
- All employees have a name
- All employees have a salary (default 10000)
- Some employees have phone numbers, if not we use NULLs

```
CREATE TABLE Employee
(
    Name CHAR(50)
        NOT NULL,
    Salary INT
        DEFAULT 10000
        NOT NULL,
    Phone CHAR(15)
        NULL
);
```

# SQL Support

- SQL allows you to insert NULLs
- You can also check for NULLs

```
INSERT INTO Employee  
VALUES  
( 'John' ,12000 ,NULL);
```

```
SELECT Name FROM  
Employee WHERE  
Phone IS NULL;
```

```
UPDATE Employee  
SET Phone = NULL  
WHERE Name = 'Mark';
```

```
SELECT Name FROM  
Employee WHERE  
Phone IS NOT NULL;
```

# Which Method to Use?

- Method to use usually depends on scenario
  - Default values should not be used when they might be confused with 'real' values
  - Splitting tables shouldn't be used too much or you'll have lots of tables
  - NULLs can be (and often are) used where the other approaches seem inappropriate
  - You don't have to always use the same method - you can mix and match as needed

# Example

- For an online store we have a variety of products - books, CDs, and DVDs
  - All items have a title, price, and id (their catalogue number)
  - Any item might have an extra shipping cost, but some don't
- There is also some data specific to each type
  - Books must have an author and might have a publisher
  - CDs must have an artist
  - DVDs might have a producer or director

# Example

- We could put all the data in one table

Items

| ID | Title | Price | Shipping | Author | Publisher | Artist | Producer | Director |
|----|-------|-------|----------|--------|-----------|--------|----------|----------|
|----|-------|-------|----------|--------|-----------|--------|----------|----------|

- Every row will have missing information
- **We are storing three types of thing in one table**
- Many additional issues that will be covered next lectures

# Example

- It is probably best to split the three types into separate tables
- We'll have a main Items table
- Also have Books, CDs, and DVDs tables with FKs to the Items table

Items

| ID | Title | Price | Shipping |
|----|-------|-------|----------|
|----|-------|-------|----------|

Books

| ID | Author | Publisher |
|----|--------|-----------|
|----|--------|-----------|

CDs

| ID | Artist |
|----|--------|
|----|--------|

DVDs

| ID | Producer | Director |
|----|----------|----------|
|----|----------|----------|

# Example

- Each of these tables might still have some missing information
- Shipping cost in items could have a default value of 0
  - This should not disrupt computations
  - If no value is given, shipping is free
- Other columns could allow NULLs
  - Publisher, director, and producer are all optional
  - It is unlikely we'll ever use them in computation

# Next Lecture

- Normalisation
  - Data Redundancy
  - Functional Dependencies
  - Normal Forms
  - First, Second and Third Normal Forms
- Further reading
  - The Manga Guide to Databases, Chapter 3
  - Database Systems, Chapter 14