

Database Security

G51DBS Database Systems

Jason Atkin

jaa@cs.nott.ac.uk

This Lecture

- General Database Security
- Privileges
 - Granting
 - Revoking
- Views
- SQL Injection Attacks
- Further Reading
 - The Manga Guide to Databases, Chapter 5
 - Database Systems, Chapter 20

Database Security

- Database security is about controlling access to information
 - Some information should be available freely
 - Other information should only be available to certain people or groups
- Many aspects to consider for security
 - Physical security
 - OS/Network security
 - Encryption and passwords
 - DBMS security
- This lecture will focus on DBMS security
 - The other elements are not really G51DBS concerns

DBMS Security Support

- DBMSs can provide some security
 - Each user has an account, username and password
 - These are used to identify a user and control their access to information
- The DBMS verifies password and checks a user's permissions when they try to:
 - Retrieve data
 - Modify data
 - Modify the database structure

Permissions and Privilege

- SQL uses privileges to control access to tables and other database objects. E.g.
 - SELECT privilege *
 - INSERT privilege *
 - UPDATE privilege *
 - CREATE privilege

** - these can apply to specific columns*
- In MySQL there are actually 30 distinct privileges
- The owner (creator) of a database has all privileges on all objects in the database, and can grant these to others
- The owner (creator) of an object has all privileges on that object and can pass them on to others

Privileges in SQL

- Will use PostgreSQL (not MySQL) – more ‘standard’

```
GRANT <privileges>  
ON <object>  
TO <users>  
[WITH GRANT OPTION]
```

- **WITH GRANT OPTION** means that the users can pass their privileges on to others
- “If WITH GRANT OPTION is specified, the recipient of the privilege may in turn grant it to others...”
<http://www.postgresql.org/docs/8.2/static/sql-grant.html>

- **<privileges>** is a list of
SELECT (<columns>),
INSERT (<columns>),
DELETE,
UPDATE (<columns>),
or simply **ALL**

- **<users>** is a list of user
- **<object>** is the name of a table or view

Privileges Examples

GRANT ALL ON Employee

TO Manager WITH GRANT OPTION;

- The user 'Manager' can do **anything** to the Employee table, **and** can allow other users to do the same (by using GRANT statements)

GRANT SELECT, UPDATE(Salary) ON Employee

TO Finance;

- The user 'Finance' can **view the entire Employee table**, and can **change Salary values**, but cannot change other values or pass on their privilege

Removing Privileges

- If you want to remove a privilege or grant option you have granted you use:

REVOKE

```
[GRANT OPTION FOR]
<privileges>
ON <object>
FROM <users>
[CASCADE | RESTRICT];
```

- For example:

REVOKE


```
UPDATE(Salary)
ON Employee
FROM Finance
```

REVOKE

```
GRANT OPTION FOR
ALL PRIVILEGES
FROM Manager
```



Removing Privileges

- Example
 - 'Admin' grants **ALL** privileges to 'Manager', and **SELECT** to 'Finance' with grant option
 - 'Manager' grants **ALL** to Personnel
 - 'Finance' grants **SELECT** to Personnel




Removing Privileges

- Manager' revokes ALL from 'Personnel'
 - 'Personnel' still has SELECT privileges from 'Finance'



Removing Privileges

- Manager' revokes ALL from 'Personnel'
 - 'Personnel' still has SELECT privileges from 'Finance'
- 'Admin' revokes SELECT from 'Finance' (with cascade)
 - Personnel also loses SELECT, since it only had it from 'Finance'



Views

- Privileges work at the level of tables
 - You **can** restrict access by column
 - You **cannot** restrict access by row
- Views, along with privileges, allow for customised access
 - i.e. Views allow you to limit access to only certain rows or columns
- Views provide 'virtual' tables
 - A view is the result of a **SELECT** statement which is treated like a table
 - You can **SELECT** from (and sometimes **UPDATE** etc) views just like tables

Creating Views

```
CREATE VIEW <name>  
AS  
<select statement>;
```

- <name> is the name of the new view
- <select statement> is a query that returns the rows and columns of the view

- Example
 - We want each university tutor to be able to see marks of only those students they actually teach
 - We will assume our database is structured with Student, Enrolment, Tutors and Module tables similar to those seen in previous lectures

View Example : the tables

Student

| | | | |
|------------|---------------|--------------|--------------|
| sID | sFirst | sLast | sYear |
|------------|---------------|--------------|--------------|

Enrolment

| | | | |
|------------|--------------|--------------|-------------------|
| sID | mCode | eMark | eYearTaken |
|------------|--------------|--------------|-------------------|

Module

| | | |
|--------------|---------------|-----------------|
| mCode | mTitle | mCredits |
|--------------|---------------|-----------------|

Tutors

| | |
|------------|------------|
| IID | sID |
|------------|------------|

Lecturers

| | | |
|------------|--------------|--------------|
| IID | IName | IDept |
|------------|--------------|--------------|

View Example

```
CREATE VIEW TuteeMarks AS
SELECT sID, sFirst, sLast, mCode, eMark
FROM Student INNER JOIN Enrolment USING(sID)
           INNER JOIN Module USING (mCode)
WHERE sID IN (SELECT sID FROM Tutors
              WHERE lID = CURRENT_USER);
```

Grant permissions to access this to every tutor:

```
GRANT SELECT ON TuteeMarks TO 'tutorname'@'%' ;
```

Note: CURRENT_USER is the current mysql user. This is called USER in Oracle. % means 0 or more characters. Only 'need' to quote *user_name* or *host_name* if it contains control characters. In Oracle you can grant to PUBLIC. MYSQL does not allow wildcard usernames.

Live example... (coursework 2 table)

- CREATE SQL SECURITY INVOKER VIEW DList
AS SELECT description FROM Description;
- REVOKE SELECT ON DList FROM jaa;
- GRANT SELECT ON DList to jaa;
- SELECT * FROM DList;
- REVOKE SELECT ON DList FROM jaa;
- DROP VIEW DList;
- SELECT * FROM DList;

Database Integrity

- **Database Security**

- Database security makes sure that the user is authorised to access information
- Beyond security, checks should be made that user mistakes are detected and prevented

- **Database Integrity**

- Ensures that authorised users only input consistent data into the database
- Usually consists of a series of constraints and assertions on data

Database Integrity

- Integrity constraints come in a number of forms:
 - **CHECK** can be used in a column definition, e.g.:
 - `gender CHAR NOT NULL
CHECK(gender IN ('M', 'F'))`
 - **DOMAINs** can be used to create custom types, e.g.:
 - `CREATE DOMAIN gendertype AS CHAR
DEFAULT 'F' CHECK (VALUE IN ('M', 'F'));`
 - **ASSERTIONS** can be used to validate checks across multiple tables/columns, assert something must always be true
- Note: check the syntax for your database. Oracle supports CHECK constraints. MySQL seems to accept the syntax and often ignore it. MySQL can emulate these with triggers – see CREATE TRIGGER ...

Connections to a DBMS

- A major concern with database security should be when your application connects to the DBMS
 - The user doesn't connect to the DBMS, the application does
 - **This often happens with elevated privileges**
 - If the application isn't well secured, it could provide a conduit for malicious code

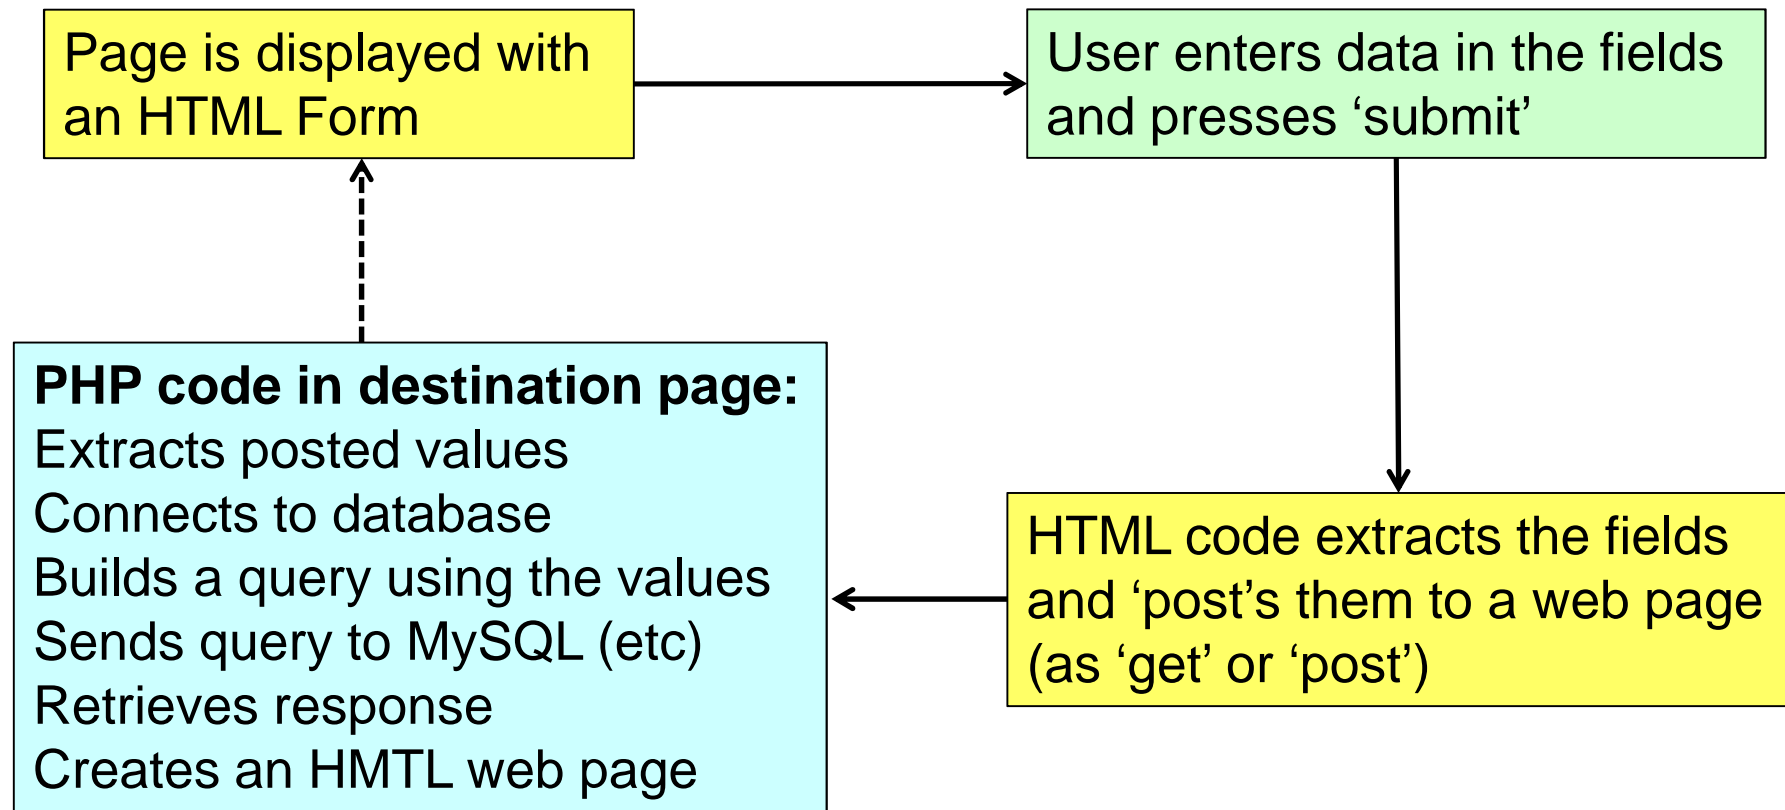
SQL Injection Attacks

An SQL Injection attack is an **exploit** where a user is able to insert **malicious** code into an SQL query resulting in an entirely new query!

WARNING!!!

- Data Protection Act 1998, Section 55(1):
A person must not knowingly or recklessly, without the consent of the data controller obtain or disclose personal data or the information contained in personal data.
- **Do not try this on a website you do not own**
- *"I was just seeing if it would work"* is not a valid defence

Usual web page / PHP interaction



e.g. see Lecture 14, <http://www.cs.nott.ac.uk/~jaa/dbs/htmlform.php>
or <http://www.cs.nott.ac.uk/~jaa/dbs/exercise4/dbdemo.php>

SQL Injection Attacks : The Cause

- If a page sends a search string in a parameter called 'searchfor', PHP code could do the following:

```
$search = $_POST['searchfor']
```

```
$query = "SELECT * FROM Products WHERE  
        pName LIKE '%" . $search . "%'";
```

- In many cases the posted value came from something which the user typed into a web page
- If user can pass malicious information, this information may be combined with regular SQL queries
- The resulting query may have a very different effect

SQL Injection Attacks : The Mistake

- An application or website is vulnerable to an injection attack if the programmer hasn't added code to check for special characters in the input:
 - ' represents the beginning or end of a string
 - ; represents the end of a command
 - /*...*/ represent comments
 - -- represents a comment for the rest of a line

SQL Injection Attacks : Example

- e.g. User login webpage, requests user ID & password
- Passed from form to PHP, as 'id' and 'pass'
- ID later used for a query, which should take the form:

```
SELECT uPass FROM Users WHERE uID = 'Jason';
```

- Example PHP code using the variable 'id':

```
$query = "SELECT uPass FROM Users WHERE uID  
= '$_POST['id']'";
```

```
$result = mysql_query($query);
```

```
$row = mysql_fetch_row($result);
```

```
$pass = row['uPass']; (Compare vs input)
```

- Password then compared with the provided password field

SQL Injection Attacks

- If the user enters *Name*, the command becomes:

```
SELECT uPass FROM Users  
WHERE uID = 'Name' ;
```

- But what about if the user entered

```
' ;DROP TABLE Users;--
```


as their **name**?

SQL Injection Attacks

- The website programmer intended to execute a single SQL query:

```
SELECT uPass FROM Users WHERE uID = 'Name'
```

String
Concatenation

```
SELECT uPass FROM Users WHERE uID = 'Name'
```

SQL Injection Attacks

- With the malicious code inserted, the meaning of the SQL changes into two queries and a comment:

```
SELECT uPass FROM Users WHERE uID = ' ';DROP TABLE Users;--'
```

String
Concatenation

```
SELECT uPass FROM Users WHERE uID = " ";DROP TABLE Users; -- '
```

SQL Injection Attacks : mysql_query()

- With the malicious code inserted, the meaning of the SQL changes into two queries and a comment:

```
SELECT uPass FROM Users WHERE uID = ' ';DROP TABLE Users;--'
```

String
Concatenation

```
SELECT uPass FROM Users WHERE uID = " ";DROP TABLE Users; -- '
```

Note: this one shouldn't actually be a problem with mysql_query() : “**mysql_query()** sends a unique query (multiple queries are not supported) to the currently active database on the server that's associated with the specified link_identifier.”²⁹

SQL Injection Attacks

- Sometimes the goal isn't sabotage, but information
- Consider an online banking system:

```
SELECT No, SortCode FROM Accounts WHERE No = '11244102'
```

String
Concatenation

```
SELECT No, SortCode FROM Accounts WHERE No = '11244102'
```

SQL Injection Attacks

- This attack is aimed at listing all accounts at a bank. The SQL becomes a single, altered query:

```
SELECT No, SortCode FROM Accounts WHERE No = '1' OR 'a' = 'a'
```

String
Concatenation

```
SELECT No, SortCode FROM Accounts WHERE No = '1' OR 'a' = 'a'
```

Particularly effective with weakly typed languages like PHP

Defending Against Injection Attacks

- Defending against SQL injection attacks is not difficult, but many people still don't
- There are many ways to improve security
- You should be doing most of these all of the time when a user inputs values that will be used in an SQL statement
- **Summary: Don't trust that all users will do what you expect them to do**

1. Restrict DBMS Access Privileges

- Assuming an SQL injection attack is successful, a user will have access to tables based on the privileges of the account that the application used to connect to the DBMS
- GRANT an application or website the minimum possible access to the database
- Do not allow DROP, DELETE etc unless absolutely necessary
- Use Views to hide as much as possible

2. Encrypt Sensitive Data

- Storing sensitive data inside your database can always lead to problems with security
- If in doubt, encrypt sensitive information so that if any breaches occur, damage is minimal
- Another reason to encrypt data is that the majority of commercial security breaches are '*inside jobs*' by trusted employees
- **Never** store unencrypted passwords, although many shops still do this 😞

3a. Validate Input

- Always validate the input values
- Arguably the most important consideration when creating a database or application that handles user input is to validate it
- Filter any escape characters and check the length of the input against expected sizes
- Checking input length should be standard practice. This applies to programming in general, as it also avoids buffer overflow attacks

3b. Validate Input

- ***Always*** escape special characters. All languages that execute SQL strings should allow this, e.g.:

```
$username = mysql_real_escape_string($input);  
$query = "SELECT * FROM Users  
        WHERE uID = '" . $username . "'";  
$result = mysql_query($query);
```

- `mysql_real_escape_string()` will escape any special characters, like ', with \
- **You should do this with all input variables**

4. Check Input Types

- In weakly typed languages, check that the user is providing you with a type you'd expect
- For example, if you expect the ID to be an int, make sure it is. In PHP:

```
if ( !is_int($_POST['userid']) )  
{  
    // ID is not an integer  
}
```

5. Stored Procedures

- Some DBMSs allow you to store procedures for use over and over again
- Procedures you might store are SELECTs, INSERTS etc, or other procedural code
- This adds another level of abstraction between the user and the tables
- If necessary, a stored procedure can access tables that are restricted to the rest of the application

6. Use Generic Error Messages

- While it might seem helpful to output informative error messages, this actually supplies users with far too much information
- For example, if your SQL query fails, **do not show the user `mysql_error()`**, instead output:
A system error has occurred. We apologise for the inconvenience.
- You can **log** the error **privately** for administrative purposes

7a. Use Parameterised Input

- Parameterised input essentially means that user input is passed to the database as parameters, not as part of the SQL string
- This makes injection attacks extremely difficult
- Not all DBMSs / Languages support this
- In PHP, you need to use PHP Data Objects (PDO) to provide parameterisation
- Reference:
<http://php.net/manual/en/book.pdo.php>

7b. Use PDO prepare()

- PHP Data Objects provides a Prepare() function:

```
$conn = new PDO(
    'mysql:host=localhost;dbname=test',$user, $pass );
$stmt = $conn->prepare(
    'SELECT * FROM Users WHERE uName = :name' );
$stmt->bindValue( ':name', $_POST['username'] );
$stmt->execute( );
```

- Rather than building up a string for your SQL using the posted variable...
- The statement is pre-compiled during prepare
- A malicious parameter may still be passed to the query, but it is simply used as the variable, not as part of the statement

Next Lecture

- Transactions
 - ACID Properties
 - COMMIT and ROLLBACK
- Recovery
 - System and Media Failures
- Concurrency

Transactions and Recovery

G51DBS Database Systems

Jason Atkin

jaa@cs.nott.ac.uk

This Lecture

- Transactions
 - ACID Properties
 - COMMIT and ROLLBACK
- Recovery
 - System and Media Failures
- Concurrency
- Further reading
 - The Manga Guide to Databases, Chapter 5
 - Database Systems, Chapter 22

Transactions

- A transaction is an action, or a series of actions, carried out by a single user or an application program, which reads or updates the contents of a database.
- All database access by users is thought of in terms of transactions

Transactions

- A transaction is a 'logical unit of work' on a database
 - Each transaction does something on the database
 - No part of it alone achieves anything useful or of interest
- Transactions are the unit of recovery, consistency and integrity
- ACID properties
 - Atomicity
 - Consistency
 - Isolation
 - Durability

Atomicity

- Transactions are atomic
 - Conceptually do not have component parts
 - In reality a transaction may include numerous read, write and other operations
- Transactions **cannot** be executed partially
 - Either performed entirely, or not at all
 - It should not be **detectable** that they interleave with another transaction

Consistency

- Transactions take the database from one **consistent** state to another
- Consistency isn't guaranteed part-way through a transaction
 - Because of atomicity, this won't be a problem
- Enforced by the DBMS, and application programmers also have some responsibility

Isolation

- All transactions execute independently of one another
- The effects of a transaction are invisible to other transactions until it has been completed
- Enforced by the scheduler

Durability

- Once a transaction has completed, its changes are made permanent
- If the database system crashes, completed transactions must remain complete
- Enforced by the recovery manager

Transaction Example

- Transfer £50 from bank account A to account B

Read(A)
A = A - 50
Write(A)
Read(B)
B = B + 50
Write(B)

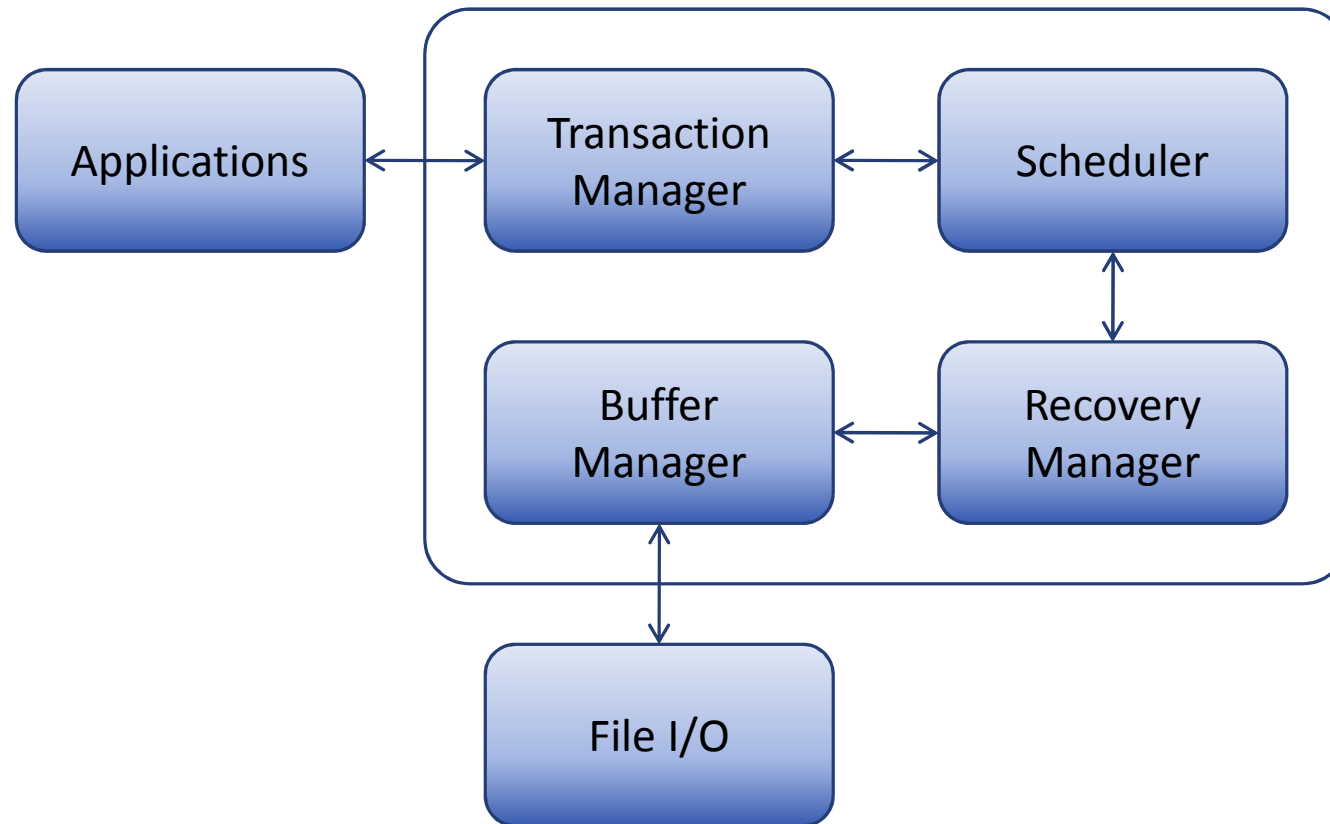
} **Transaction**

- **Atomicity** – Shouldn't take money from A without giving it to B
- **Consistency** – Money isn't lost or gained overall
- **Isolation** – Other queries shouldn't see A or B change until completion
- **Durability** – The money does not return to A, even after a system crash

Transaction Subsystem

- The transaction subsystem enforces the ACID properties
 - Schedules the operations of all transactions
 - Uses COMMIT and ROLLBACK to ensure atomicity
 - Locks and/or timestamps are used to ensure consistency and isolation (next lectures)
 - A log is kept to ensure durability

Transaction Subsystem



Database Systems, Connolly & Begg, p574

COMMIT and ROLLBACK

- COMMIT is used to signal the successful end of a transaction
 - Any changes that have been made to the database should be made permanent
 - These changes are now available to other transactions
- ROLLBACK is used to signal the unsuccessful end of a transaction
 - Any changes that have been made to the database should be undone
 - It is now as if the transaction never happened, it can now be reattempted if necessary

Recovery

- Transactions must be durable, but some failures will be unavoidable
 - System crashes
 - Power failures
 - Disk crashes
 - User mistakes
 - Sabotage
 - etc
- Prevention is better than a cure
 - Reliable OS
 - Security
 - UPS and surge protectors
 - RAID arrays
- Can't protect against everything, system recovery will be necessary

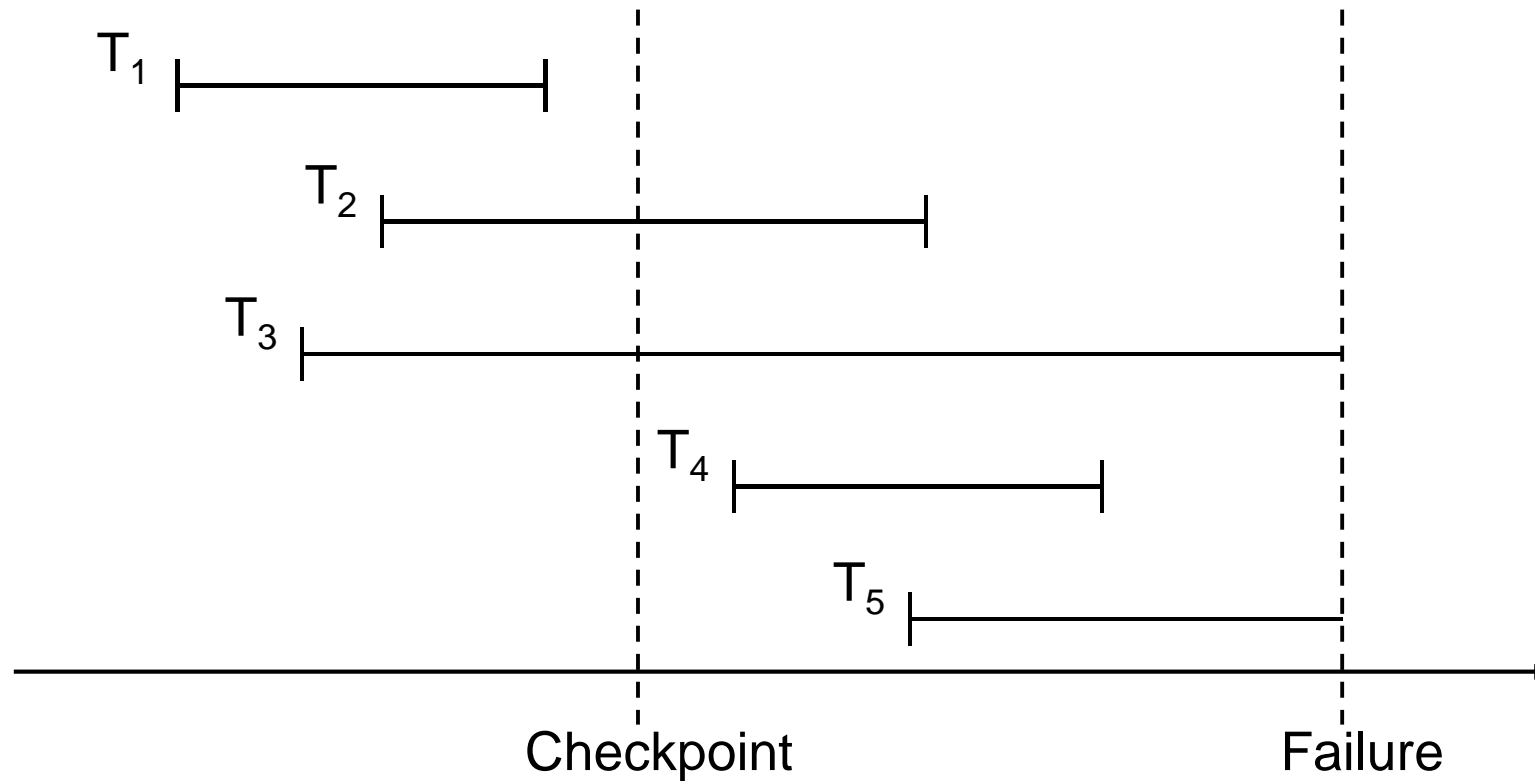
The Transaction Log

- The transaction log records details of all transactions
 - Any changes the transaction makes to the database
 - How to undo these changes
 - When transactions complete and how
- **The log is stored on disk, not in memory**
 - If the system crashes, the log is preserved
- **Write ahead log rule**
 - The entry in the log must be made before COMMIT processing can complete

System Failures

- A system failure affects all running transactions
 - Software crash
 - Power failure
- The physical media (disks) are not damaged
- Things in memory are lost, things on the disk are kept
- At various times a DBMS takes a checkpoint
 - All transactions are written to disk
 - A record is made (on disk) of all transactions that are currently running
- Until that time, changes could be in memory but not on disk

Transaction Timeline



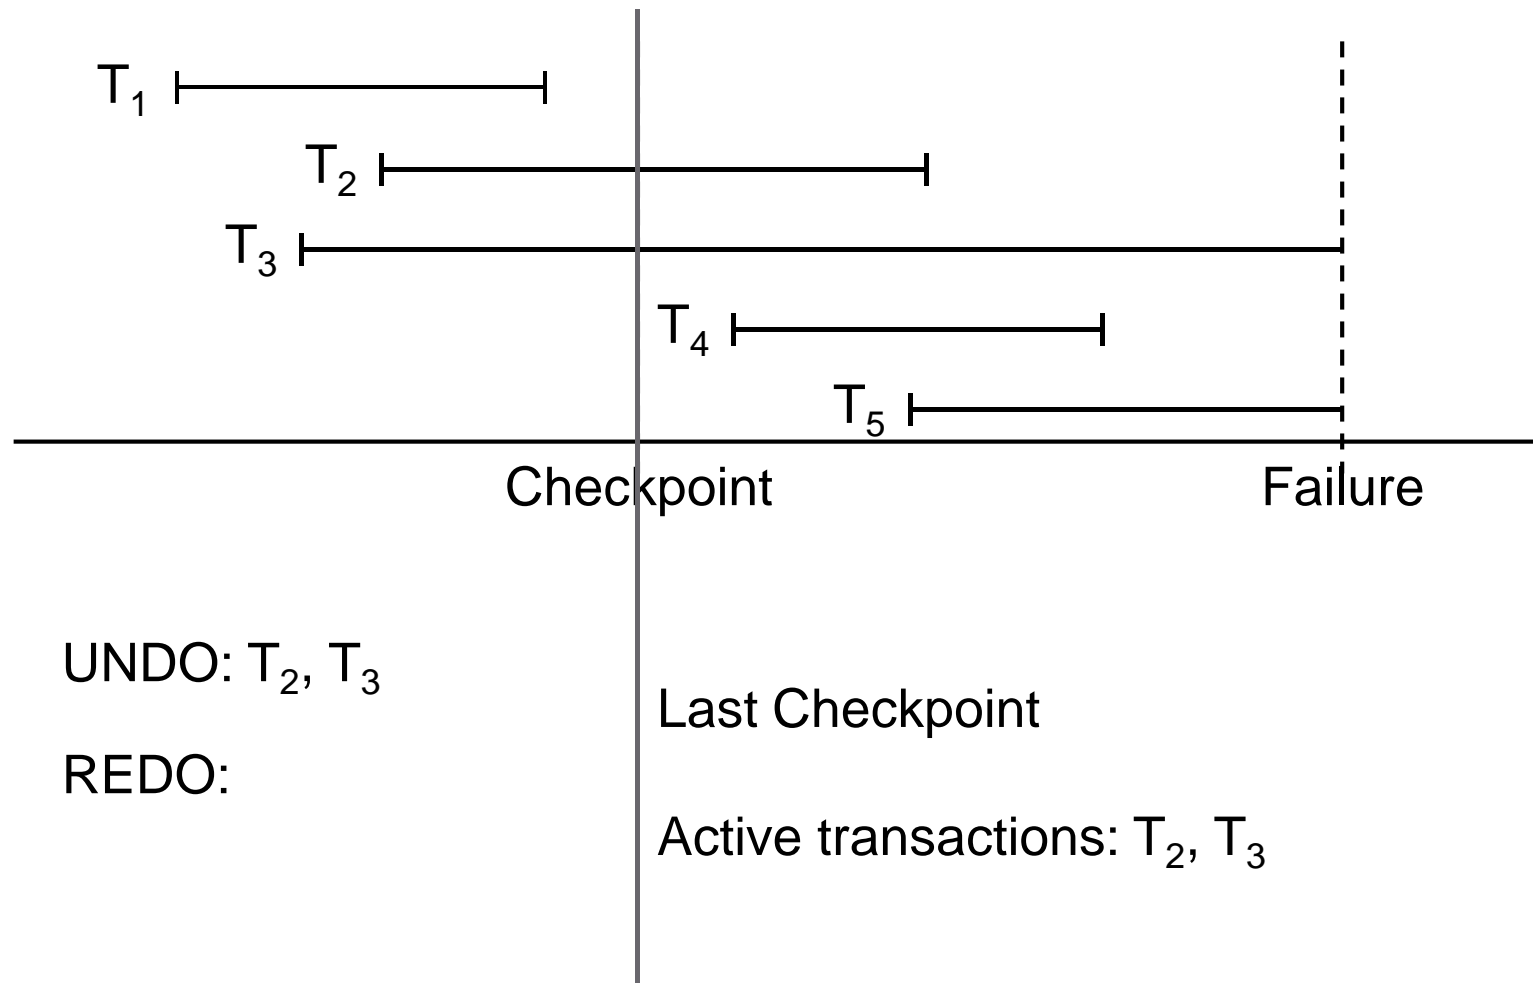
System Recovery

- Any transaction that was running at the time of failure needs to be undone and possibly restarted
- Any transactions that committed since the last checkpoint need to be redone
- Transactions of type T_1 (completed before check point) need no recovery
- Transactions of type T_3 or T_5 (uncompleted) need to be undone
- Transactions of type T_2 or T_4 (completed SINCE check point) need to be redone

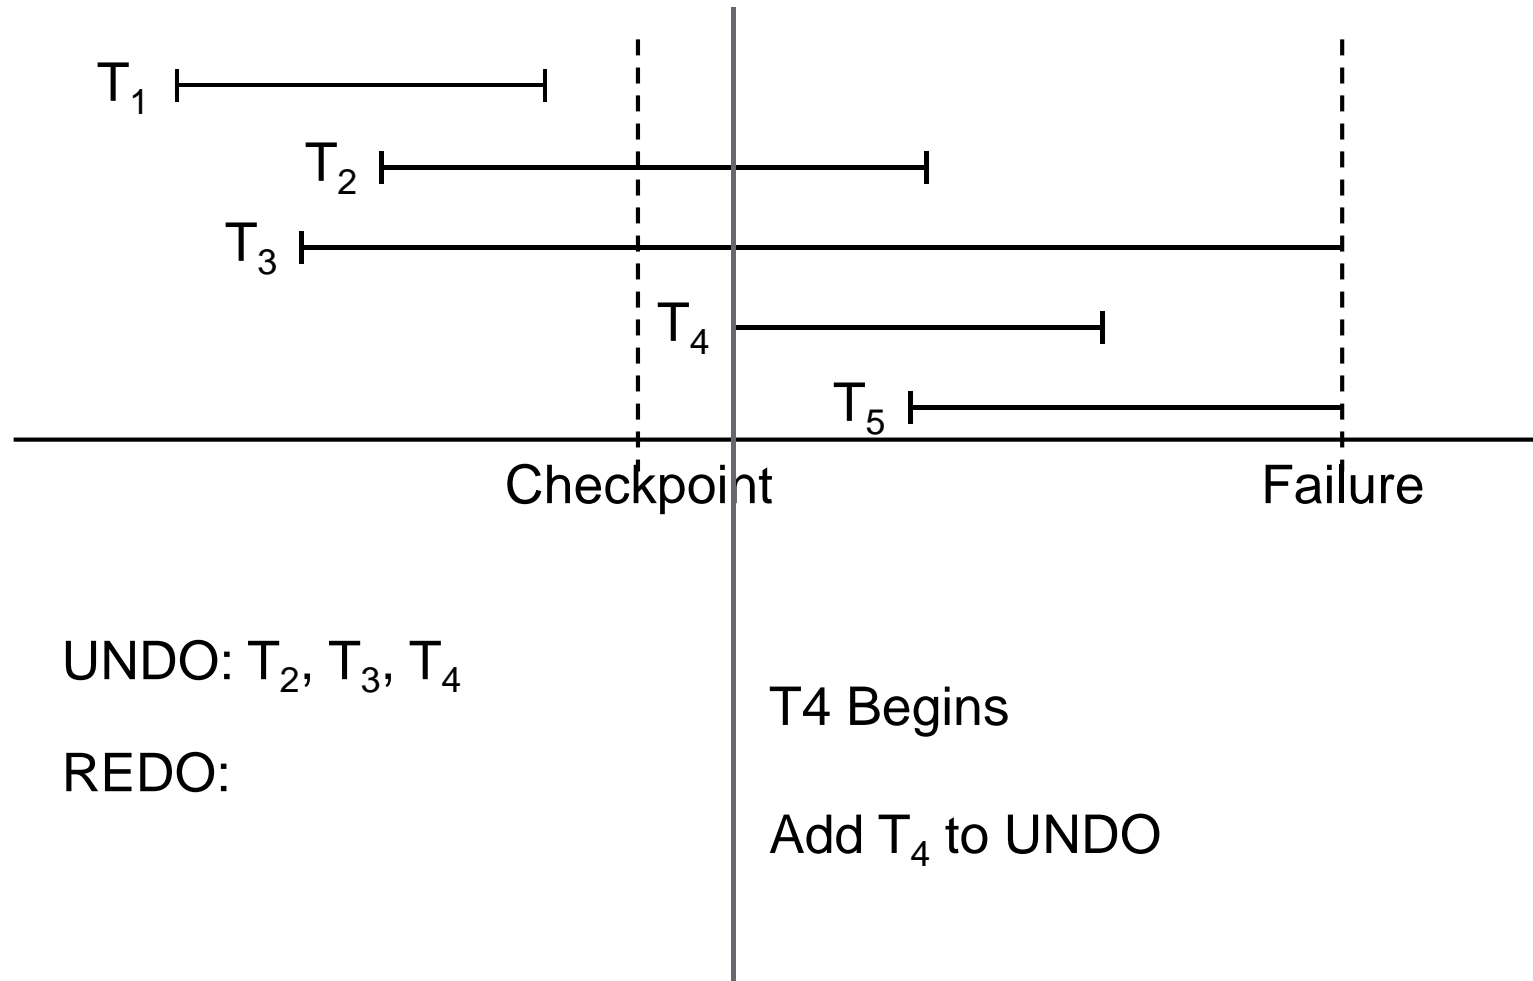
Transaction Recovery

- Create two lists of transactions: UNDO and REDO
 - UNDO – all transactions running at the last checkpoint
 - REDO – empty
- For every entry in the log since the last checkpoint, until the failure:
 1. If a BEGIN TRANSACTION entry is found for T:
 - Add T to UNDO
 2. If a COMMIT entry is found for T:
 - Remove T From UNDO
 - Add T to REDO (it had finished)

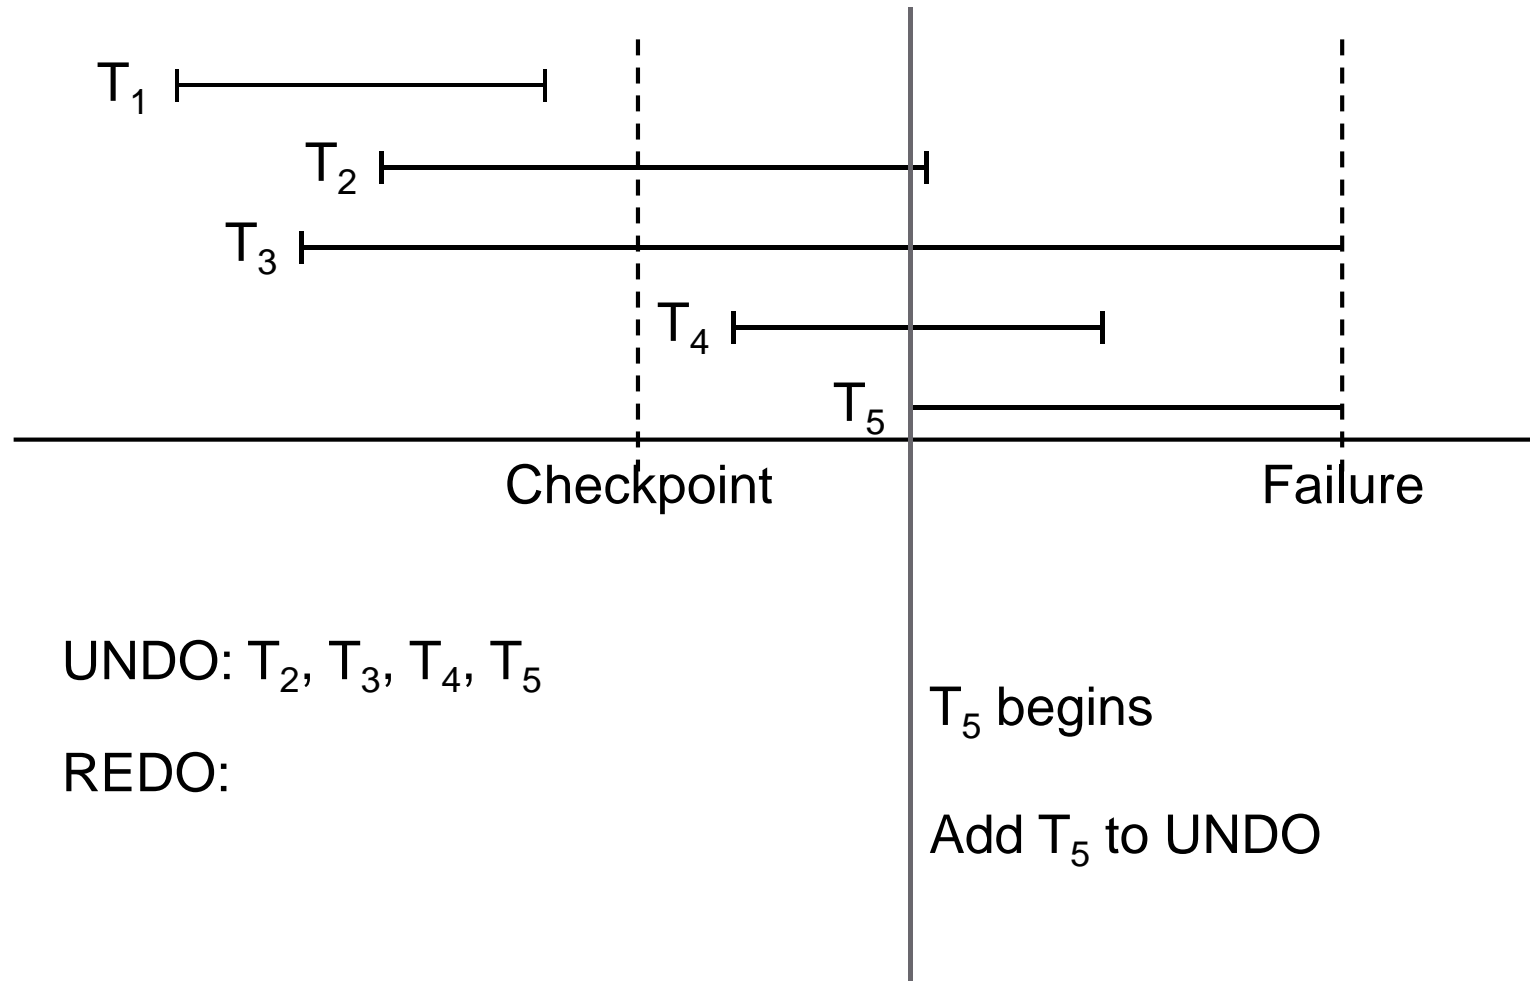
Transaction Recovery



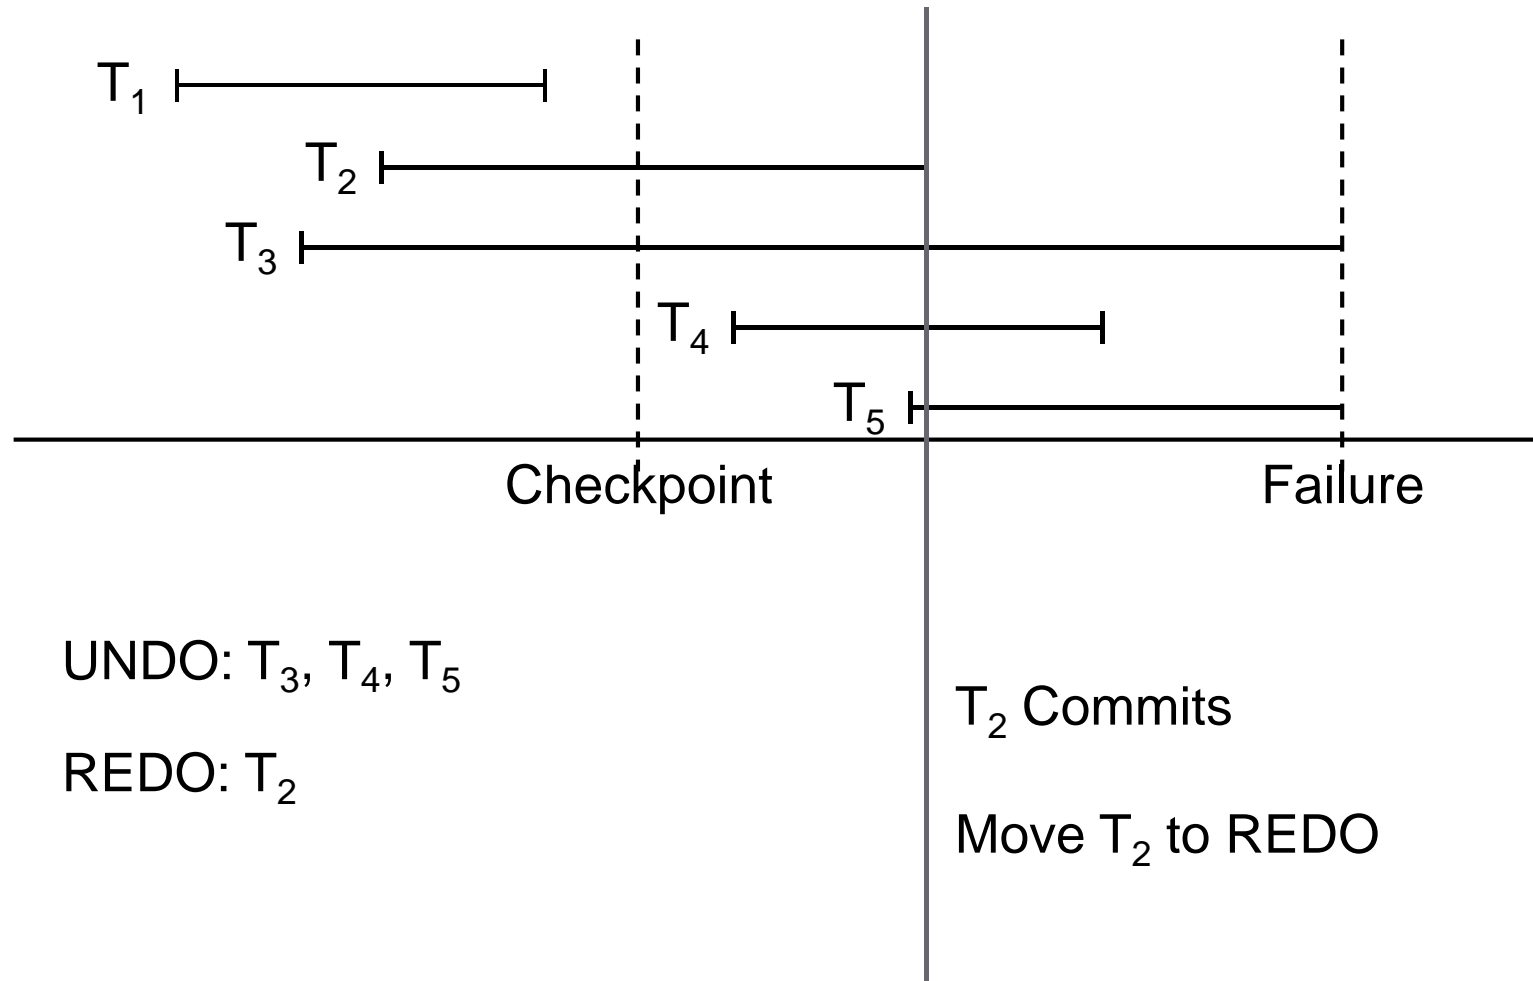
Transaction Recovery



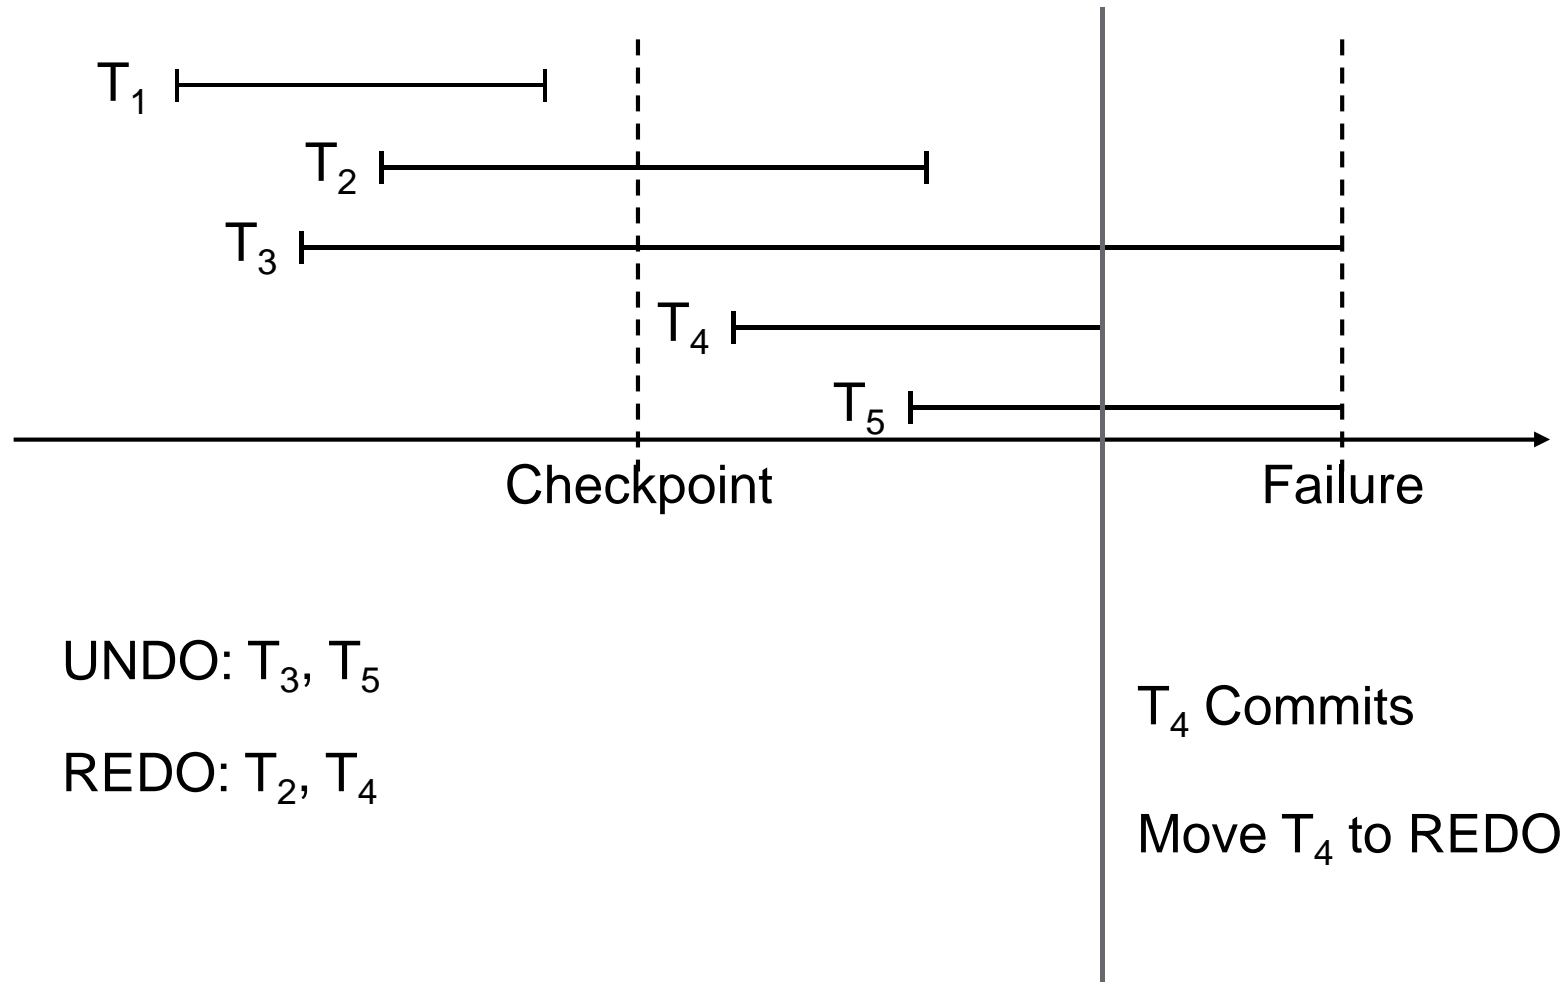
Transaction Recovery



Transaction Recovery



Transaction Recovery



Forwards and Backwards

- Backwards recovery - ROLLBACK
 - We need to undo some transactions
 - Working backwards through the log we undo every operation by any transaction on the UNDO list
 - This returns the database to a consistent state – although with some uncompleted transactions (those on the redo list)
- Forwards recovery - ROLLFORWARD
 - Some transactions need to be redone
 - Working forwards through the log we redo any operation by a transaction on the REDO list
 - This brings the database up to date

Media Failures

- System failures are not too severe
 - Only information since the last checkpoint is affected
 - This can be recovered from the transaction log
- Media failures (e.g. Disk failure) are more serious
 - The stored data is damaged
 - The transaction log itself may be damaged

Backups

- Backups are necessary to recover from media failure
 - The transaction log and entire database is written to secondary storage
 - Very time consuming, often requires downtime
- Backup frequency
 - Frequent enough that little information is lost
 - Not so frequent as to cause problems
 - Every night is a common compromise

Recovery from Media Failure

1. Restore the database from the last backup
 2. Use the transaction log to redo any changes made since the last backup
- If the transaction log is damaged you can't do step 2
 - Store the log on a separate physical device to the database
 - This reduces the risk of losing both together

Transactions in MySQL

- Most DBMSs support transactions
- In MySQL in school only the InnoDB engine supports transactions
- There are other engines which support this, that are not installed, like Falcon
- On the school servers, autocommit is set so that every command is instantly committed
- This is very slow and inefficient
- And does not make it easy to undo changes
- You can turn autocommit off with

```
SET autocommit = 0 | 1;
```

Managing Transactions

- In MySQL, a transaction is executed in the following way:

```
BEGIN | START TRANSACTION;  
INSERT INTO table VALUES (...);  
SELECT col1, col2 FROM table;  
UPDATE table SET col1 = col2 + 3;  
DROP TABLE table;  
COMMIT | ROLLBACK;
```

(| *optional*)

Managing Transactions

- In PHP, you can send off these commands with `mysql_query`:

```
mysql_query( 'BEGIN' );  
mysql_query( '...' );  
if (some test)  
{  
    mysql_query( 'COMMIT' );  
}  
else  
{  
    mysql_query( 'ROLLBACK' );  
}
```


Managing Transactions

- In general, this approach is far superior to autocommit. Remember, however:
 - If your transaction locks a table, all other transactions will have to wait
 - So COMMIT as soon as possible
 - MyISAM and most engines ignore commands like ROLLBACK. So use InnoDB if you need transaction support
 - Subqueries are good when using autocommit to avoid outdated information

Concurrency

- Large databases are used by many people
 - Many transactions are to be run on the database
 - It is helpful to run these simultaneously
 - Still need to preserve isolation
- If we don't allow for concurrency then transactions are run sequentially
 - Have a queue of transactions
 - Easy to preserve atomicity and isolation
 - Long transactions (e.g. backups) will delay others

Concurrency Problems

- In order to run two or more concurrent transactions, their operations must be interleaved
- Each transaction gets a share of the computing time
- This can lead to several problems
 - Lost updates
 - Uncommitted updates
 - Incorrect updates
- **All arise when isolation is broken**
 - i.e. we want a way to avoid this

Lost Update

| T1 | T2 |
|------------------------|------------------------|
| Read(X) $X = X - 5$ | |
| | Read(X) $X = X + 5$ |
| Write(X) | Write(X) |
| COMMIT | COMMIT |

- T1 and T2 both read X, both modify it, then both write it out
 - The net effect of both transactions should be no change to X
 - Only T2's change is seen however

Uncommitted Update

| T1 | T2 |
|------------------------------------|--|
| Read(X) $X = X - 5$ Write(X) | |
| ROLLBACK | Read(X) $X = X + 5$ Write(X) COMMIT |

- T2 sees the change to X made by T1, but T1 is then rolled back
 - The change made by T1 is rolled back
 - i.e. the change would not actually be made
 - X goes back to what it was prior to the transaction
 - But T2 already used it
 - It should be as if that change never happened

Inconsistent Analysis

| T1 | T2 |
|------------------------------------|--|
| Read(X) $X = X - 5$ Write(X) | |
| Read(Y) $Y = Y + 5$ Write(Y) | Read(X) Read(Y) $\text{Sum} = X + Y$ |

- T1 doesn't change the sum of X and Y, but T2 records a change
 - T1 consists of two parts - take 5 from X then add 5 to Y
 - T2 sees the effect of the first change, but not the second

Resolutions for these problems

- We need a way to stop multiple transactions reading/updating the same data if it will cause a problem
- We need to know which things can be done concurrently (e.g. read + read) and which things cannot be done (e.g. read + write, write + write)
- We need a way to either stop something from changing (Locking, lecture 17), or to know it was changed and 'roll back' to restart if we hit a problem (Timestamping, lecture 18)
 - Each approach has advantages and disadvantages

This Lecture in Exams

Define a transaction in the context of database management

Explain how a DBMS uses a transaction log to recover from a system failure using ROLLBACK and ROLLFORWARD

Explain the difference between a system failure and a media failure

Next Lecture

- Concurrency
 - Locks and Resources
 - Deadlock
- Serialisability
 - Schedules of transactions
 - Serial and serialisable schedules
- Further reading
 - The Manga Guide to Databases, Chapter 5
 - Database Systems, Chapter 22